

# ✈ Flight Fare Prediction using Machine Learning

## Problem Statement

Flight fare prediction involves applying data science and machine learning to estimate the cost of airline tickets based on various flight details. These include airline, travel date, source and destination cities, number of stops, and duration. By using statistical modeling on historical flight data, machine learning models can identify pricing patterns and forecast ticket fares. This enables smarter travel planning and supports dynamic pricing strategies.

### Target:

- Design a predictive model using machine learning algorithms to forecast the ticket fare of a flight based on its features – this is a regression problem.
- Design an analytical model to identify and analyze the most significant factors influencing flight fares, helping airlines and travelers understand pricing trends – this is an insight-driven exploratory analysis task.

### Dataset Description:

Column Name	Description
Unnamed: 0	Index/serial number, not relevant for analysis
airline	Airline name (e.g., Vistara, IndiGo, etc.)
flight	Flight number (e.g., UK-706)
source_city	City from which the flight departs
departure_time	Time of the day when flight departs (Morning, Evening, etc.)
stops	Number of stops (non-stop, one stop, two or more)
arrival_time	Time of the day when flight arrives
destination_city	City where flight lands
class	Travel class ( Economy , Business )
duration	Total duration of flight (in hours, float)
days_left	Days left for departure from the date of booking
price	<b>Target Variable</b> – Flight fare (in INR)

### Importing Necessary Libraries

```
In [34]: #Data Manipulation Libraries
import pandas as pd
import numpy as np

# Data Visualization Libraries
import seaborn as sns
import matplotlib.pyplot as plt

# Warning Libraries
import warnings
warnings.filterwarnings("ignore")

# Data preprocessing Libraries
from sklearn.preprocessing import OneHotEncoder
OHE = OneHotEncoder()

from sklearn.preprocessing import OrdinalEncoder
ORE = OrdinalEncoder()

# Model_selection Libraries
from sklearn.model_selection import train_test_split

# Machine Learning Model
from sklearn.ensemble import RandomForestRegressor
LAR = RandomForestRegressor()
# RFR = RandomForestRegressor()
from sklearn.linear_model import Lasso
# LAR = Lasso(alpha=0.1)

# Evaluation Metrics
from sklearn.metrics import mean_squared_error
```

```
from sklearn.metrics import r2_score
```

```
# GUI Libraries
```

```
import tkinter as tk
```

```
from tkinter import ttk, messagebox
```

## Importing Dataset

```
In [35]: airline_dataset = pd.read_csv("AirlinesDataset.csv")
```

## Exploratory Data Analysis

Dataset before elimination of unnecessary columns

```
In [36]: airline_dataset
```

```
Out[36]:
```

	Unnamed: 0	airline	flight	source_city	departure_time	stops	arrival_time	destination_city	class	duration	days_le
0	0	SpiceJet	SG-8709	Delhi	Evening	zero	Night	Mumbai	Economy	2.17	
1	1	SpiceJet	SG-8157	Delhi	Early_Morning	zero	Morning	Mumbai	Economy	2.33	
2	2	AirAsia	I5-764	Delhi	Early_Morning	zero	Early_Morning	Mumbai	Economy	2.17	
3	3	Vistara	UK-995	Delhi	Morning	zero	Afternoon	Mumbai	Economy	2.25	
4	4	Vistara	UK-963	Delhi	Morning	zero	Morning	Mumbai	Economy	2.33	
...	...	...	...	...	...	...	...	...	...	...	...
300148	300148	Vistara	UK-822	Chennai	Morning	one	Evening	Hyderabad	Business	10.08	4
300149	300149	Vistara	UK-826	Chennai	Afternoon	one	Night	Hyderabad	Business	10.42	4
300150	300150	Vistara	UK-832	Chennai	Early_Morning	one	Night	Hyderabad	Business	13.83	4
300151	300151	Vistara	UK-828	Chennai	Early_Morning	one	Evening	Hyderabad	Business	10.00	4
300152	300152	Vistara	UK-822	Chennai	Morning	one	Evening	Hyderabad	Business	10.08	4

300153 rows × 12 columns

\*The columns named 'Unnamed: 0' and 'flight' are not required for training this machine learning model, so they have been eliminated.

Eliminating Unnecessary columns from the dataset

```
In [37]: airline_dataset = airline_dataset.drop("Unnamed: 0",axis=1)
airline_dataset = airline_dataset.drop("flight",axis=1)
airline_dataset = airline_dataset.drop("duration",axis=1)
airline_dataset = airline_dataset.drop("arrival_time",axis=1)
```

Checking first 10 entries of the dataset -

```
In [38]: airline_dataset.head(10)
```

```
Out[38]:
```

	airline	source_city	departure_time	stops	destination_city	class	days_left	price
0	SpiceJet	Delhi	Evening	zero	Mumbai	Economy	1	5953
1	SpiceJet	Delhi	Early_Morning	zero	Mumbai	Economy	1	5953
2	AirAsia	Delhi	Early_Morning	zero	Mumbai	Economy	1	5956
3	Vistara	Delhi	Morning	zero	Mumbai	Economy	1	5955
4	Vistara	Delhi	Morning	zero	Mumbai	Economy	1	5955
5	Vistara	Delhi	Morning	zero	Mumbai	Economy	1	5955
6	Vistara	Delhi	Morning	zero	Mumbai	Economy	1	6060
7	Vistara	Delhi	Afternoon	zero	Mumbai	Economy	1	6060
8	GO_FIRST	Delhi	Early_Morning	zero	Mumbai	Economy	1	5954
9	GO_FIRST	Delhi	Afternoon	zero	Mumbai	Economy	1	5954

Checking last 10 entries of the dataset -

```
In [39]: airline_dataset.tail(10)
```

```
Out[39]:
```

	airline	source_city	departure_time	stops	destination_city	class	days_left	price
300143	Air_India	Chennai	Early_Morning	one	Hyderabad	Business	49	51345
300144	Air_India	Chennai	Evening	one	Hyderabad	Business	49	51345
300145	Air_India	Chennai	Morning	one	Hyderabad	Business	49	51345
300146	Air_India	Chennai	Early_Morning	one	Hyderabad	Business	49	51345
300147	Air_India	Chennai	Early_Morning	one	Hyderabad	Business	49	68739
300148	Vistara	Chennai	Morning	one	Hyderabad	Business	49	69265
300149	Vistara	Chennai	Afternoon	one	Hyderabad	Business	49	77105
300150	Vistara	Chennai	Early_Morning	one	Hyderabad	Business	49	79099
300151	Vistara	Chennai	Early_Morning	one	Hyderabad	Business	49	81585
300152	Vistara	Chennai	Morning	one	Hyderabad	Business	49	81585

Shape of the Dataset -

```
In [40]: print("Shape of the above dataset:",airline_dataset.shape)
print("Numbers of columns:",12)
print("Numbers of Rows:",300153)
```

Shape of the above dataset: (300153, 8)

Numbers of columns: 12

Numbers of Rows: 300153

Columns Names -

```
In [41]: airline_dataset.columns
```

```
Out[41]: Index(['airline', 'source_city', 'departure_time', 'stops', 'destination_city',
               'class', 'days_left', 'price'],
              dtype='object')
```

Checking for null or missing values -

```
In [42]: airline_dataset.isnull().sum()
```

```
Out[42]: airline          0
source_city          0
departure_time       0
stops                0
destination_city     0
class                0
days_left           0
price                0
dtype: int64
```

Obtaining some information about the dataset -

```
In [43]: airline_dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 300153 entries, 0 to 300152
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   airline          300153 non-null object
1   source_city      300153 non-null object
2   departure_time   300153 non-null object
3   stops           300153 non-null object
4   destination_city 300153 non-null object
5   class            300153 non-null object
6   days_left        300153 non-null int64
7   price            300153 non-null int64
dtypes: int64(2), object(6)
memory usage: 18.3+ MB
```

Datatypes present in the dataset -

```
In [44]: airline_dataset.dtypes
```

```
Out[44]: airline            object
source_city               object
departure_time            object
stops                    object
destination_city          object
class                    object
days_left                int64
price                    int64
dtype: object
```

Separating categorical columns from dataset

```
In [45]: categorical_Columns = []
for i in airline_dataset.columns:
    if airline_dataset[i].dtype == object:
        categorical_Columns.append(i)
categorical_Columns
```

```
Out[45]: ['airline',
'source_city',
'departure_time',
'stops',
'destination_city',
'class']
```

Separating numerical columns from dataset

```
In [46]: numerical_Columns = []
for i in airline_dataset.columns:
    if airline_dataset[i].dtype != object:
        numerical_Columns.append(i)
numerical_Columns
```

```
Out[46]: ['days_left', 'price']
```

Merging both columns for Checking distinct values

```
In [47]: combined_columns = categorical_Columns + numerical_Columns
combined_columns
```

```
Out[47]: ['airline',
'source_city',
'departure_time',
'stops',
'destination_city',
'class',
'days_left',
'price']
```

Checking distinct values in each column

```
In [48]: value_counts_dict = {}
for column in airline_dataset.columns:
    value_counts_dict[column] = airline_dataset[column].value_counts()
for column, counts in value_counts_dict.items():
    print("-" * (len(column)+4) )
    print("|",column,"|")
    print("-" * (len(column)+4) )
    print(counts)
    print("=====")
```

```
-----
| airline |
-----
```

```
airline
Vistara      127859
Air_India    80892
Indigo       43120
GO_FIRST     23173
AirAsia      16098
SpiceJet     9011
Name: count, dtype: int64
=====
```

```
-----
| source_city |
-----
```

```
source_city
Delhi        61343
Mumbai       60896
Bangalore    52061
Kolkata      46347
Hyderabad    40806
```

```

Chennai      38700
Name: count, dtype: int64
=====
-----
| departure_time |
-----
departure_time
Morning      71146
Early_Morning 66790
Evening      65102
Night        48015
Afternoon    47794
Late_Night   1306
Name: count, dtype: int64
=====
-----
| stops |
-----
stops
one          250863
zero         36004
two_or_more  13286
Name: count, dtype: int64
=====
-----
| destination_city |
-----
destination_city
Mumbai      59097
Delhi       57360
Bangalore   51068
Kolkata     49534
Hyderabad   42726
Chennai     40368
Name: count, dtype: int64
=====
-----
| class |
-----
class
Economy     206666
Business    93487
Name: count, dtype: int64
=====
-----
| days_left |
-----
days_left
25      6633
18      6602
39      6593
32      6585
26      6573
24      6542
19      6537
31      6534
33      6532
40      6531
41      6525
28      6522
38      6512
20      6502
30      6501
42      6497
22      6494
36      6490
21      6479
37      6476
43      6472
44      6436
17      6419
11      6417
34      6412
13      6404
23      6401
29      6397
12      6381
27      6360
14      6349
15      6340
45      6314
35      6291
16      6272

```

```

46    6160
49    6154
48    6078
47    6069
10    5822
8     5767
6     5740
7     5703
9     5665
5     5392
4     5077
3     4248
2     4026
1     1927

```

Name: count, dtype: int64

```
=====
```

```

-----
| price |
-----

```

```

price
54608    1445
2339     1442
54684    1390
60978    1383
60508    1230

```

```

...
10767     1
10946     1
12206     1
12972     1
5607      1

```

Name: count, Length: 12157, dtype: int64

```
=====
```

Encoding the columns for better accuracy of the model

```
In [49]: airline_dataset_encoded = airline_dataset.copy()
```

\*we applied One-Hot Encoding to the columns['airline', 'source\_city', 'departure\_time', 'arrival\_time', 'destination\_city'] because they are categorical features, and machine learning models require numerical input to perform calculations.

```
In [50]: OHE = OneHotEncoder(drop='first', sparse_output=False, handle_unknown='ignore')
onehot_encoding_columns = ['airline', 'source_city', 'departure_time', 'destination_city']
columns_to_encode = airline_dataset_encoded[onehot_encoding_columns]
encoded_array = OHE.fit_transform(columns_to_encode)
encoded_df = pd.DataFrame(encoded_array, columns=OHE.get_feature_names_out(onehot_encoding_columns), index=airline_dataset_encoded.index)
airline_dataset_encoded.drop(columns=onehot_encoding_columns, inplace=True)
airline_dataset_encoded = pd.concat([airline_dataset_encoded, encoded_df], axis=1)
airline_dataset_encoded.head(10)
```

```
Out[50]:
```

	stops	class	days_left	price	airline_Air_India	airline_GO_FIRST	airline_Indigo	airline_SpiceJet	airline_Vistara	source_city
0	zero	Economy	1	5953	0.0	0.0	0.0	1.0	0.0	
1	zero	Economy	1	5953	0.0	0.0	0.0	1.0	0.0	
2	zero	Economy	1	5956	0.0	0.0	0.0	0.0	0.0	
3	zero	Economy	1	5955	0.0	0.0	0.0	0.0	1.0	
4	zero	Economy	1	5955	0.0	0.0	0.0	0.0	1.0	
5	zero	Economy	1	5955	0.0	0.0	0.0	0.0	1.0	
6	zero	Economy	1	6060	0.0	0.0	0.0	0.0	1.0	
7	zero	Economy	1	6060	0.0	0.0	0.0	0.0	1.0	
8	zero	Economy	1	5954	0.0	1.0	0.0	0.0	0.0	
9	zero	Economy	1	5954	0.0	1.0	0.0	0.0	0.0	

10 rows × 24 columns

```
◀ | ▶
```

\*We applied ordinal encoding because the number of stops has a natural increasing order that affects fare. This helps the model understand that more stops usually imply longer/cheaper flights.

```
In [51]: stops_order = [['zero', 'one', 'two_or_more']]
ORE = OrdinalEncoder(categories=stops_order)
airline_dataset_encoded['stops_encoded'] = ORE.fit_transform(airline_dataset_encoded[['stops']])
airline_dataset_encoded['stops'] = airline_dataset_encoded['stops_encoded']
airline_dataset_encoded.drop(columns='stops_encoded', inplace=True)
```

\*We used ordinal encoding because flight classes (Economy < Business) have a clear price hierarchy. Encoding preserves this relationship so the model can learn its impact on fare.

```
In [52]: class_order = [['Economy', 'Business']]
ORE_class = OrdinalEncoder(categories=class_order)
airline_dataset_encoded['class_encoded'] = ORE_class.fit_transform(airline_dataset_encoded[['class']])
airline_dataset_encoded['class'] = airline_dataset_encoded['class_encoded']
airline_dataset_encoded.drop(columns='class_encoded', inplace=True)
```


Random sample of the dataset just to ensure that the encoding is accurate

```
In [53]: airline_dataset_encoded.sample(20)
```

Out[53]:

	stops	class	days_left	price	airline_Air_India	airline_GO_FIRST	airline_Indigo	airline_SpiceJet	airline_Vistara	source_city_Chennai
140908	1.0	0.0	39	4056	0.0	0.0	0.0	0.0	0.0	
213140	1.0	1.0	16	57017	0.0	0.0	0.0	0.0	1.0	
107132	2.0	0.0	32	7106	0.0	0.0	0.0	0.0	1.0	
58056	2.0	0.0	30	10006	0.0	0.0	0.0	0.0	1.0	
90207	1.0	0.0	32	2723	0.0	0.0	0.0	0.0	0.0	
64553	1.0	0.0	19	4772	0.0	0.0	1.0	0.0	0.0	
15502	1.0	0.0	28	5040	0.0	0.0	1.0	0.0	0.0	
70196	1.0	0.0	49	4977	1.0	0.0	0.0	0.0	0.0	
166593	1.0	0.0	13	4453	0.0	0.0	1.0	0.0	0.0	
234326	1.0	1.0	4	65517	0.0	0.0	0.0	0.0	1.0	
145441	2.0	0.0	27	7123	1.0	0.0	0.0	0.0	0.0	
209992	0.0	1.0	32	36712	0.0	0.0	0.0	0.0	1.0	
33430	1.0	0.0	36	6172	0.0	0.0	0.0	0.0	1.0	
22603	1.0	0.0	16	5686	1.0	0.0	0.0	0.0	0.0	
205510	1.0	0.0	36	3979	1.0	0.0	0.0	0.0	0.0	
136140	1.0	0.0	6	13524	1.0	0.0	0.0	0.0	0.0	
210486	1.0	1.0	36	50969	1.0	0.0	0.0	0.0	0.0	
119245	2.0	0.0	44	9847	0.0	0.0	0.0	0.0	1.0	
276517	0.0	1.0	36	24122	1.0	0.0	0.0	0.0	0.0	
219124	0.0	1.0	1	38470	0.0	0.0	0.0	0.0	1.0	

20 rows × 24 columns



Splitting features and target

```
In [54]: X = airline_dataset_encoded.drop("price",axis=1)
Y = airline_dataset_encoded["price"]
```

Displaying the Features -

```
In [55]: X
```

Out[55]:

	stops	class	days_left	airline_Air_India	airline_GO_FIRST	airline_Indigo	airline_SpiceJet	airline_Vistara	source_city_Chennai
0	0.0	0.0	1	0.0	0.0	0.0	1.0	0.0	
1	0.0	0.0	1	0.0	0.0	0.0	1.0	0.0	
2	0.0	0.0	1	0.0	0.0	0.0	0.0	0.0	
3	0.0	0.0	1	0.0	0.0	0.0	0.0	1.0	
4	0.0	0.0	1	0.0	0.0	0.0	0.0	1.0	
...	...	...	...	...	...	...	...	...	
300148	1.0	1.0	49	0.0	0.0	0.0	0.0	1.0	
300149	1.0	1.0	49	0.0	0.0	0.0	0.0	1.0	
300150	1.0	1.0	49	0.0	0.0	0.0	0.0	1.0	
300151	1.0	1.0	49	0.0	0.0	0.0	0.0	1.0	
300152	1.0	1.0	49	0.0	0.0	0.0	0.0	1.0	

300153 rows × 23 columns



Displaying the Target -

```
In [56]: Y
```

```
Out[56]: 0          5953
         1          5953
         2          5956
         3          5955
         4          5955
         ...
        300148      69265
        300149      77105
        300150      79099
        300151      81585
        300152      81585
        Name: price, Length: 300153, dtype: int64
```

Splitting Training and testing data

```
In [57]: X_train , X_test , Y_train , Y_test = train_test_split(X,Y , test_size= 0.2 ,random_state=14)
```

Training the model -

```
In [58]: LAR.fit(X_train, Y_train)
```

```
Out[58]: ▼ RandomForestRegressor ⓘ ?
         RandomForestRegressor()
```

Predicting the test data -

```
In [59]: y_pred = LAR.predict(X_test)
```

Evaluating the accuracy of the model -

```
In [60]: mse = mean_squared_error(Y_test, y_pred)
         r2 = r2_score(Y_test, y_pred)


         print("Mean Squared Error:", mse)
         print("R² Score:", r2)
         print("Predictions:",y_pred)
```

```
Mean Squared Error: 20425400.763711303
R² Score: 0.9604541327637137
Predictions: [ 8032.7136039  62495.6386014  44633.01467857 ... 54608.
              9110.571      7167.47266638]
```

```
In [61]: plt.figure(figsize=(8,5))
         sns.scatterplot(x=Y_test, y=y_pred, alpha=0.3)
         plt.xlabel("Actual Price")
         plt.ylabel("Predicted Price")
         plt.title("Actual vs Predicted Flight Fare")
         plt.show()
```



Actual vs Predicted Flight Fare



Predictive System -

```
In [62]: airline_categories = ['AirAsia', 'AirIndia', 'GoFirst', 'Indigo', 'SpiceJet', 'Vistara']
city_categories = ['Bangalore', 'Chennai', 'Delhi', 'Hyderabad', 'Kolkata', 'Mumbai']
time_categories = ['Early_Morning', 'Evening', 'Morning', 'Night']
stop_categories = ['zero', 'one', 'two_or_more']
class_categories = ['Economy', 'Business']

def predict_fare():
    try:
        airline = airline_var.get()
        source = source_var.get()
        departure = departure_var.get()
        stops = stops_var.get()
        dest = destination_var.get()
        f_class = class_var.get()
        days = int(days_left_var.get())

        input_df = pd.DataFrame([airline, source, departure, dest],
                                columns=['airline', 'source_city', 'departure_time', 'destination_city'])
        encoded_array = OHE.transform(input_df)
        encoded_df = pd.DataFrame(encoded_array,
                                   columns=OHE.get_feature_names_out(['airline', 'source_city', 'departure_time',
                                                                        'destination_city']),
                                   index=[0])

        stops_encoded = ORE.transform([[stops]])[0][0]
        class_encoded = ORE_class.transform([[f_class]])[0][0]

        final_input = pd.DataFrame([stops_encoded, class_encoded, days], columns=['stops', 'class', 'days_left'])
        final_input = pd.concat([final_input, encoded_df], axis=1)

        missing = set(X.columns) - set(final_input.columns)
        for col in missing:
            final_input[col] = 0
        final_input = final_input[X.columns]

        result = LAR.predict(final_input)[0]
        result_label.config(text=f"➡ Estimated Flight Fare: ₹{round(result, 2)}")

    except Exception as e:
        messagebox.showerror("Error", f"Something went wrong: {e}")
```

Graphical User Interface -

```
In [63]: root = tk.Tk()
root.title("Flight Fare Predictor")

tk.Label(root, text="Select Flight Details", font=("Arial", 14)).grid(row=0, column=0, columnspan=2, pady=10)
tk.Label(root, text="Airline").grid(row=1, column=0)
```

```
airline_var = ttk.Combobox(root, values=airline_categories)
airline_var.grid(row=1, column=1)

tk.Label(root, text="Source City").grid(row=2, column=0)
source_var = ttk.Combobox(root, values=city_categories)
source_var.grid(row=2, column=1)

tk.Label(root, text="Departure Time").grid(row=3, column=0)
departure_var = ttk.Combobox(root, values=time_categories)
departure_var.grid(row=3, column=1)

tk.Label(root, text="Stops").grid(row=4, column=0)
stops_var = ttk.Combobox(root, values=stop_categories)
stops_var.grid(row=4, column=1)

tk.Label(root, text="Destination City").grid(row=5, column=0)
destination_var = ttk.Combobox(root, values=city_categories)
destination_var.grid(row=5, column=1)

tk.Label(root, text="Class").grid(row=6, column=0)
class_var = ttk.Combobox(root, values=class_categories)
class_var.grid(row=6, column=1)

tk.Label(root, text="Days Left").grid(row=7, column=0)
days_left_var = tk.Spinbox(root, from_=0, to=365)
days_left_var.grid(row=7, column=1)

tk.Button(root, text="Predict Fare", command=predict_fare).grid(row=8, column=0, columnspan=2, pady=10)

result_label = tk.Label(root, text="", font=("Arial", 12), fg="green")
result_label.grid(row=9, column=0, columnspan=2)

root.mainloop()
```